

Geodetic First Approximation of Size and Timing: User's Guide

Brendan Crowell ^{*1} and Ben Baker ^{†2}

¹Pacific Northwest Seismic Network

²Instrumental Software Technologies, Inc.

September 21, 2016

Contents

1	Introduction	3
2	Installation	4
2.1	Prerequisites	4
2.1.1	Compilers	4
2.1.2	CMake	4
2.1.3	LAPACK and BLAS	4
2.1.4	libGeographic	5
2.1.5	iniparser	5
2.1.6	HDF5	5
2.1.7	libxml2	5
2.1.8	ISCL	5
2.1.9	ActiveMQ	6
2.1.10	Earthworm	6
2.1.11	FFTw	6
2.1.12	Cython	6
2.1.13	cMoPaD	6
2.2	Building GFAST	6
3	GFAST Library Overview	8
3.1	GFAST Core Routines	9
3.2	Data Acquisition	10
3.3	XML Output	11
3.4	HDF5 Output	12

*crowellb@pnsn.org

†benbaker@isti.com

4 Acknowledgements	12
A License	12

1 Introduction

copy brendan’s original manual, note the amazon funding, and add brendan’s publications so people know what to cite

G-FAST (Geodetic First Approximation of Size and Timing) was developed by Brendan Crowell as a Python based geodetic earthquake early warning (EEW) module. Funding from Amazon has allowed PNSN to develop a compiled-language production variant of G-FAST for integration into the USGS’s next generation ShakeAlert EEW system. The software is released under the license described in Appendix A.

Nominally, G-FAST acquires real-time geodetic data and, when triggered by a ShakeAlert message, activates three geodetic co-seismic inversion suites: a peak ground displacement (PGD) inversion, a centroid moment tensor (CMT) inversion, and a CMT driven slip or finite fault inversion. The PGD module provides a quick and robust estimate of magnitude. The CMT inversion provides a robust estimate of magnitude, depth, and moment tensor components. The moment-tensor derived fault planes and depth are then passed onto the slip inversion which resolves the CMT fault plane ambiguity while providing an additional magnitude estimate. The PGD and slip model are then passed onto the ShakeAlert decision module as the event unfolds.

As in the Python implementation, G-FAST still acquires real-time data from the Pacific Northwest Geodetic Array (PANGA) and, in the EEW context, is triggered on ShakeAlert messages. The key differences are now that an Earthworm server is responsible for the import of PANGA data and the ShakeAlert communication is directly implemented via ActiveMQ. Additionally, G-FAST has been separated into multiple C-libraries to encourage reuse by groups outside of EEW.

2 Installation

This section outlines the GFAST installation procedure. This consists of two primary steps. The first step is to install the dependencies. The second step is to then build GFAST and link to the dependencies.

2.1 Prerequisites

This section lists the requisite, optional, and workaround libraries for building GFAST.

2.1.1 Compilers

The following libraries and GFAST can be compiled with a C, C++, and Fortran compiler. It is recommended the compilers have complete OpenMP-4 support for improved code vectorization. OpenMP-4 support is realized in GCCv5 and above or Clang 3.9 and above though GFAST has been successfully built with GCCv4.4. If using GCCv4 then it is recommended one set the “-Wno-unknown-pragmas” C compiler flag.

2.1.2 CMake

CMake is used because GFAST is decomposed into many smaller libraries with varying dependencies. This dependency ambiguity is implicitly handled by this automatic makefile generator. In addition to generation of makefiles also CMake provides for automatic generation and execution of unit tests. As with any makefile creation step the correct specification of dependencies can be quite painful but, upon successful generation of a CMakeCache.txt file this activity need not be repeated and subsequent pulls from the git repository should build without hassle. CMake is likely available through a package manager or can be obtained from <https://cmake.org/>. GFAST has been successfully built with Version 2.6 and above.

2.1.3 LAPACK and BLAS

GFAST ultimately aims to solve matrix-vector equations of the form $\|G\mathbf{m} = \mathbf{d}\|_2$ (LAPACK) or estimate data by computing matrix-vector multiplies of the form $\mathbf{u} = G\mathbf{m}$ (BLAS). Moreover, profiling shows that computation of the least-squares problem, particularly in the finite fault inversion, is the greatest contributor to program execution time. Necessarily, one can easily achieve better performance by using high-performance (commercial) LAPACK and BLAS libraries. Consequently, it is recommended one follow the strategy of

1. Check for an existing vendor LAPACK/BLAS implementation.
2. If using an Intel processor obtain Intel’s Math Kernel Library (MKL) which is freely available at <https://software.intel.com/sites/campaigns/nest/>. Also, note that MKL contains an implementation of FFTw and may reduce the number of dependencies.

3. Check a package manager for a pre-built LAPACK and BLAS which are distributed with a modified BSD license.
4. TODO: look for OpenBLAS then, if not found, ATLAS in the CMakeModules if using a package manager's lapacke + blas.
5. At last resort obtain and build LAPACK and BLAS from <http://www.netlib.org/lapack/> which are distributed with a modified BSD license.

2.1.4 libGeographic

This is a library used by GFAST's coordinate tools unit tests. It may be available through a package manager or, if need be, obtained from <http://geographiclib.sourceforge.net/>. It is distributed under the MIT license.

2.1.5 iniparser

This is the utility which parses the GFAST parameter file. It is available from <https://github.com/ndevilla/iniparser> and is distributed under the MIT license.

2.1.6 HDF5

The HDF5 self-describing file format is used for high-performance and portable data archival and data playback. HDF5 can be obtained from a package manager or from <https://www.hdfgroup.org/HDF5/>. GFAST appears to work with HDF5 1.18.6. If static linking to HDF5 then one must additionally obtain the zlib compression library. A zlib implementation is available in Intel's Performance Primitives available at <https://software.intel.com/sites/campaigns/nest/>. Zlib can also be obtained with a package manager or from <http://www.zlib.net>. Notice that there is no data compression in GFAST disk writes and this step can be avoided by linking to the shared HDF5 library. If choosing between IPP and a zlib then note that IPP improves the performance of ISCL. Zlib and HDF5 both are distributed with permissive software licenses.

2.1.7 libxml2

The ShakeAlert and QuakeML data products are created with libxml2. If using GCC then it is very likely you already have libxml2. Otherwise, it can be obtained from <http://xmlsoft.org/>. libxml2 is distributed under the MIT license.

2.1.8 ISCL

Brendan's initial GFAST implementation was written in Python and made extensive use of Numerical Python. Somewhat fortuitously ISTI had been developing a library targeted at expediting the conversion of NumPy laden Python scripts to C via the ISTI Scientific Computing Library (ISCL). ISCL is distributed under the Apache-2 license and freely available from <https://github.com/bakerb845/libiscl>.

2.1.9 ActiveMQ

The ShakeAlert earthquake early warning system sends and receives alerts with the ActiveMQ messaging system. If one is not using GFAST in earthquake early warning then ActiveMQ is not necessary. ActiveMQ additionally depends on libcrypto and libssl. ActiveMQ is distributed under the Apache 2 license is available at <http://activemq.apache.org/cms/download.html>.

2.1.10 Earthworm

The real-time GPS data are currently incorporated into GFAST via Earthworm which is available at <http://earthworm.isti.com/trac/earthworm/>. If not using the real-time system then Earthworm is not necessary. Earthworm requires two additional libraries, RabbitMQ-c and Jansson for import of the GeoJSON which contains the GPS precise-point-position data. RabbitMQ-c is available at <https://github.com/alanxz/rabbitmq-c> and is distributed with an MIT license. Jansson is available at <https://github.com/akheron/jansson> and is distributed with a permissive license.

2.1.11 FFTw

The requirement for this library will be dictated by your linking policy. If you require static linking *and* are not using Intel MKL then to successfully build GFAST you would have to obtain FFTw from a package manager or obtain it from <http://www.fftw.org/>. Notice however that no Fourier transforms are computed in GFAST. Additionally, if static linking without MKL then FFTw's GPL-2 copyleft will impact GFAST's internal license. To avoid this complication it is recommended one use MKL or link to the dynamic ISCL library.

2.1.12 Cython

A developmental Python interface is provided through Cython which converts Python to C and is available at <http://cython.org/> or through a package manager. This is not required and at the present time minimally supported.

2.1.13 cMoPaD

The moment tensor decompositions are accomplished through a C implementation of MoPaD. This is distributed with GFAST. Because the C implementation is derivative work it must assume MoPaD's original LPGL3 license. This means that one can safely link to the shared MoPaD library without modifying the GFAST license. This is reflected in the target link library of the `src/CMakeLists.txt` file.

2.2 Building GFAST

To increase the likelihood that G-FAST can be compiled on multiple platforms (e.g. Linux, Windows, OSX) and ease the burden on developers to modify cumbersome makefiles G-FAST adopts an automatic makefile generation utility in CMake. Naively, one can descend into G-FAST's `src` directory and type

cmake .

The CMakeModules will attempt to find the requisite libraries and may be moderately successful if items are installed in ‘standard’ locations. Alternatively, if one would rather specify paths to custom directories cmake can be used by

```
cmake ./ -DCMAKE_BUILD_TYPE=DEBUG \  
-DCMAKE_INSTALL_PREFIX=./ \  
-DCMAKE_C_COMPILER=/usr/bin/gcc-5 \  
-DCMAKE_CXX_COMPILER=/usr/bin/c++ \  
-DCMAKE_C_FLAGS="-g3 -O2 -fopenmp -Wall" \  
-DCMAKE_CXX_FLAGS="-g3 -O2" \  
-DAPR_INCLUDE_DIR=path_to_apr \  
-DLIBAMQ_LIBRARY=path_to_libactivemq-cpp \  
-DLSSL_LIBRARY=path_to_libssl \  
-DLCRYPTO_LIBRARY=path_to_libcrypto \  
-DMKL_LIBRARY="mkl;libraries;" \  
-DIPP_LIBRARY="intel;performance;libraries" \  
-DH5_C_INCLUDE_DIR=path_to_hdf5_include \  
-DH5_LIBRARY=path_to_lib_libhdf5 \  
-DINIPARSER_INCLUDE_DIR=path_to_iniparser_src \  
-DINIPARSER_LIBRARY=path_to_libiniparser \  
-DISCL_INCLUDE_DIR=path_to_iscl_include \  
-DISCL_LIBRARY=path_to_libiscl \  
-DGEOLIB_LIBRARY=path_to_libGeographic \  
-DEW_INCLUDE_DIR=path_to_ewinclude \  
-DEW_LIBRARY=path_to_libew \  
-DLIBXML2_INCLUDE_DIR=path_to_libxml2_include \  
-DLIBXML2_LIBRARY=path_to_libxml2
```

3 GFAST Library Overview

This section will outline the GFAST libraries. Libraries are used extensively to emphasize modularity, allow users to customize GFAST to their application without incurring the cost of another GFAST application, and promote a maintainable software framework whereby software fixes can efficiently propagate to all users. The libraries to be introduced are

- Subsection 3.1: `libgfast_core` for core GFAST modeling and inversion
- add the rest

3.1 GFAST Core Routines

This section outlines the `libgfast_core` utilities which are essential to the expert earthquake-early warning driver routines. The six primary utilities are

- A properties reader for requisite GFAST modeling parameters.
- Waveform processing for computation of peak ground displacement or offsets derived from precise-point-position time series data.
- PGD inversion of peak ground displacement data computed in the waveform processing.
- CMT inversion of offset data computed in the waveform processing.
- Finite fault inversion of data computed in the waveform processing.
- Coordinate utilities for converting latitude and longitude to and from UTM¹.

¹This directory supersedes the `libGeographic` ISCL interface as `libGeographic` has unpredictable behavior when crossing UTM zones that could jeopardize program library execution in the Pacific Northwest.




Figure 1: A depiction of the current data flow. The goal is to put the data in a common, continuous, finite-sized format hereby encapsulated by the HDF5 archive. In the context of playbacks the bottom functions are not used. Hence, it is the same function that queries the HDF5 archive in playback and real-time mode and the G-FAST application developer is only exposed to one data fetcher API. The API is very simple in that G-FAST specifies two times, for example the event origin time and the current time, and the data is returned with the understanding some data may be latent or non-existent and have a predetermined value indicating its absence. For real-time scenarios the Earthworm acquisition must be configured and running prior to launching a G-FAST application. However, the GPS ring is accessed directly by a G-FAST function. In addition to a playback, a real-time application must call the G-FAST ring reader function periodically, say every second, as to refresh the HDF5 archive.

3.2 Data Acquisition

This section outlines the data acquisition strategy in real-time and playback operations. To make the real-time acquisition and playback appear integrated to the user a common HDF5 file format is selected. For playbacks, a pre-processing program can be used to generate a valid input data file² which is queried by a simple API for data in a given a time range. Likewise, real-time data is written³ to an HDF5 file with the same structure as a playback data file and, like in playback mode, can be queried by the same API as the playback for data in a given time range.

²‘File’ is a generic word in the context of this discussion. It is unnecessary and potentially undesirable from a performance standpoint that the file physically exist in the computer’s disk space. The illusion of keeping the file on disk is accomplished by instructing HDF5 to memory map the data file during the start-up phase.

³The writes can be to memory or disk.

3.3 XML Output

This section outlines the three XML data products. Two XML products, PGD and finite fault, have defined ShakeAlert schemas. The third XML product is a QuakeML summary CMT summary for potentially expediting communication between the tsunami warning centers and NEIC.

3.4 HDF5 Output

This section outlines the HDF5 summary files. HDF5 is a portable and high performance utility which makes it possible to easily archive the results⁴ of GFAST after each iteration thereby providing a snapshot of the GFAST execution and

4 Acknowledgements

This software was developed with funds from the amazon catalyst and government grants a, b, c and d.

A License

brendan, you have to email the comotion (tech transfer) and see what they want. i'm going to guess this software belongs to uw, ideally it is free (as in beer) in an academic or government setting, uw owns all modifications to the software, the software is provided as is, and you cannot hold uw liable. they may settle for a gpl-3 which would be a good compromise if the goal is to keep a company from touching it. this may have already been decided in the amazon catalyst proposal.

⁴Since results are encapsulated in data structures the HDF5 files are the data structures after each iteration.