

# Geodetic First Approximation of Size and Timing: User's Guide

Brendan Crowell <sup>\*1</sup> and Ben Baker <sup>†2</sup>

<sup>1</sup>Pacific Northwest Seismic Network

<sup>2</sup>Instrumental Software Technologies, Inc.

September 19, 2016

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>Installation</b>	<b>3</b>
2.1	Prerequisites . . . . .	3
2.1.1	Compilers . . . . .	3
2.1.2	CMake . . . . .	3
2.1.3	LAPACK and BLAS . . . . .	3
2.1.4	libGeographic . . . . .	4
2.1.5	iniparser . . . . .	4
2.1.6	HDF5 . . . . .	4
2.1.7	libxml2 . . . . .	4
2.1.8	ISCL . . . . .	4
2.1.9	ActiveMQ . . . . .	4
2.1.10	Earthworm . . . . .	5
2.1.11	FFTw . . . . .	5
2.1.12	Cython . . . . .	5
2.1.13	cMoPaD . . . . .	5
2.2	Building GFAST . . . . .	5

---

<sup>\*</sup>crowellb@pnsn.org

<sup>†</sup>benbaker@isti.com

<b>3</b>	<b>GFAST Library Overview</b>	<b>6</b>
3.1	GFAST Core Routines . . . . .	7
3.2	Data Acquisition . . . . .	8
3.3	XML Output . . . . .	9
3.4	HDF5 Output . . . . .	10

# 1 Introduction

copy brendan's original manual, note the amazon funding, and add brendan's publications so people know what to cite

## 2 Installation

This section outlines the GFAST installation procedure. This consists of two primary steps. The first step is to install the dependencies. The second step is to then build GFAST and link to the dependencies.

### 2.1 Prerequisites

This section lists the requisite, optional, and workaround libraries for building GFAST.

#### 2.1.1 Compilers

The following libraries and GFAST can be compiled with a C, C++, and Fortran compiler. It is recommended the compilers have complete OpenMP-4 support for improved code vectorization. OpenMP-4 support is realized in GCCv5 and above or Clang 3.9 and above though GFAST has been successfully built with GCCv4.4. If using GCCv4 then it is recommended one set the “-Wno-unknown-pragmas” C compiler flag.

#### 2.1.2 CMake

CMake is used because GFAST is decomposed into many smaller libraries with varying dependencies. This dependency ambiguity is implicitly handled by this automatic makefile generator. In addition to generation of makefiles also CMake provides for automatic generation and execution of unit tests. As with any makefile creation step the correct specification of dependencies can be quite painful but, upon successful generation of a CMakeCache.txt file this activity need not be repeated and subsequent pulls from the git repository should build without hassle. CMake is likely available through a package manager or can be obtained from <https://cmake.org/>.

#### 2.1.3 LAPACK and BLAS

GFAST ultimately aims to solve matrix-vector equations of the form  $\|G\mathbf{m} = \mathbf{d}\|_2$  (LAPACK) or estimate data by computing matrix-vector multiplies of the form  $\mathbf{u} = G\mathbf{m}$  (BLAS). Moreover, profiling shows that computation of the least-squares problem, particularly in the finite fault inversion, is the greatest contributor to program execution time. Necessarily, one can easily achieve better performance by using high-performance (commercial) LAPACK and BLAS libraries. Consequently, it is recommended one follow the strategy of

1. Check for an existing vendor LAPACK/BLAS implementation.
2. If using an Intel processor obtain Intel’s Math Kernel Library (MKL) which is freely available at <https://software.intel.com/sites/campaigns/nest/>. Also, note that MKL contains an implementation of FFTw and may reduce the number of dependencies.
3. Check a package manager for a pre-built LAPACK and BLAS which is distributed with a modified BSD license.

4. At last resort obtain and build LAPACK and BLAS from <http://www.netlib.org/lapack/> which is distributed with a modified BSD license.

#### 2.1.4 libGeographic

This is a library used by GFAST's coordinate tools unit tests. It may be available through a package manager or, if need be, obtained from <http://geographiclib.sourceforge.net/>. It is distributed under the MIT license.

#### 2.1.5 iniparser

This is the utility which parses the GFAST parameter file. It is available from <https://github.com/ndevilla/iniparser> and is distributed under the MIT license.

#### 2.1.6 HDF5

The HDF5 self-describing file format is used for high-performance and portable data archival and data playback. HDF5 can be obtained from a package manager or from <https://www.hdfgroup.org/HDF5/>. GFAST appears to work with HDF5 1.18.6. If static linking to HDF5 then one must additionally obtain the zlib compression library. A zlib implementation is available in Intel's Performance Primitives available at <https://software.intel.com/sites/campaigns/nest/>. Zlib can also be obtained with a package manager or from <http://www.zlib.net>. Notice that there is no data compression in GFAST disk writes and this step can be avoided by linking to the shared HDF5 library. If choosing between IPP and a zlib then note that IPP improves the performance of ISCL. Zlib and HDF5 both are distributed with permissive software licenses.

#### 2.1.7 libxml2

The ShakeAlert and QuakeML data products are created with libxml2. If using GCC then it is very likely you already have libxml2. Otherwise, it can be obtained from <http://xmlsoft.org/>. libxml2 is distributed under the MIT license.

#### 2.1.8 ISCL

Brendan's initial GFAST implementation was written in Python and made extensive use of Numerical Python. Somewhat fortuitously ISTI had been developing a library targeted at expediting the conversion of NumPy laden Python scripts to C via the ISTI Scientific Computing Library (ISCL). ISCL is distributed under the Apache-2 license and freely available from <https://github.com/bakerb845/libiscl>.

#### 2.1.9 ActiveMQ

The ShakeAlert earthquake early warning system sends and receives alerts with the ActiveMQ messaging system. If one is not using GFAST in earthquake early warning then

ActiveMQ is not necessary. ActiveMQ additionally depends on libcrypto and libssl. ActiveMQ is distributed under the Apache 2 license is available at <http://activemq.apache.org/cms/download.html>.

### 2.1.10 Earthworm

The real-time GPS data are currently incorporated into GFAST via Earthworm which is available at <http://earthworm.isti.com/trac/earthworm/>. If not using the real-time system then Earthworm is not necessary. Earthworm requires two additional libraries, RabbitMQ-c and Jansson for import of the GeoJSON which contains the GPS precise-point-position data. RabbitMQ-c is available at <https://github.com/alanxz/rabbitmq-c> and is distributed with an MIT license. Jansson is available at <https://github.com/akheron/jansson> and is distributed with a permissive license.

### 2.1.11 FFTw

The requirement for this library will be dictated by your linking policy. If you require static linking *and* are not using Intel MKL then to successfully build GFAST you would have to obtain FFTw from a package manager or obtain it from <http://www.fftw.org/>. Notice however that no Fourier transforms are computed in GFAST. Additionally, if static linking without MKL then FFTw's GPL-2 copyleft will impact GFAST's internal license. To avoid this complication it is recommended one use MKL or link to the dynamic ISCL library.

### 2.1.12 Cython

A developmental Python interface is provided through Cython which converts Python to C and is available at <http://cython.org/> or through a package manager. This is not required and at the present time minimally supported.

### 2.1.13 cMoPaD

The moment tensor decompositions are accomplished through a C implementation of MoPaD. This is distributed with GFAST. Because the C implementation is derivative work it must assume MoPaD's original LPGL3 license. This means that one can safely link to the shared MoPaD library without modifying the GFAST license. This is reflected in the target link library of the `src/CMakeLists.txt` file.

## 2.2 Building GFAST

TODO: add directions. In short type: `cmake .` then directly modify the `CMakeCache.txt`.

### 3 GFAST Library Overview

This section will outline the GFAST libraries. Libraries are used extensively to emphasize modularity, allow users to customize GFAST to their application without incurring the cost of another GFAST application, and promote a maintainable software framework whereby software fixes can efficiently propagate to all users. The libraries to be introduced are

- Subsection 3.1: `libgfast_core` for core GFAST modeling and inversion
- add the rest

### 3.1 GFAST Core Routines

This section outlines the `libgfast_core` utilities which are essential to the expert earthquake-early warning driver routines. The six primary utilities are

- A properties reader for requisite GFAST modeling parameters.
- Waveform processing for computation of peak ground displacement or offsets derived from precise-point-position time series data.
- PGD inversion of peak ground displacement data computed in the waveform processing.
- CMT inversion of offset data computed in the waveform processing.
- Finite fault inversion of data computed in the waveform processing.
- Coordinate utilities for converting latitude and longitude to and from UTM<sup>1</sup>.

---

<sup>1</sup>This directory supersedes the `libGeographic ISCL` interface as `libGeographic` has unpredictable behavior when crossing UTM zones that could jeopardize program library execution in the Pacific Northwest.



## 3.2 Data Acquisition

This section outlines the data acquisition strategy in real-time and playback operations. To make the real-time acquisition and playback appear integrated to the user a common HDF5 file format is selected. For playbacks, a pre-processing program can be used to generate a valid input data file<sup>2</sup> which is queried by a simple API for data in a given a time range. Likewise, real-time data is written<sup>3</sup> to an HDF5 file with the same structure as a playback data file and, like in playback mode, can be queried by the same API as the playback for data in a given time range.

---

<sup>2</sup>‘File’ is a generic word in the context of this discussion. It is unnecessary and potentially undesirable from a performance standpoint that the file actually exist on disk space. The illusion of keeping the file on actual disk is accomplished by instructing HDF5 to memory map the data file during the start-up phase.

<sup>3</sup>The writes can be to memory or disk.

### **3.3 XML Output**

This section outlines the three XML data products. Two XML products, PGD and finite fault, have defined shakeAlert schemas. The third XML product is a QuakeML summary CMT summary for potentially expediting communication between the tsunami warning centers and NEIC.

### 3.4 HDF5 Output

This section outlines the HDF5 summary files. HDF5 is a portable and high performance utility which makes it possible to easily archive the results<sup>4</sup> of GFAST after each iteration thereby providing a snapshot of the GFAST execution and

---

<sup>4</sup>Since results are encapsulated in data structures the HDF5 files are the data structures after each iteration.